

# text\_classify

March 3, 2020

```
In [1]: # Importing all the needed modules.
import os
from glob import glob
import matplotlib.pyplot as plt
import random
import cv2
import pandas as pd
import numpy as np
import matplotlib.gridspec as gridspec
import seaborn as sns
import zlib
import itertools
import sklearn
import itertools
import scipy
import skimage
from skimage.transform import resize
import csv
from tqdm import tqdm
import warnings
warnings.filterwarnings("ignore")
from sklearn import model_selection
from sklearn.model_selection import train_test_split, KFold, cross_val_score, StratifiedKFold
from sklearn.utils import class_weight
from sklearn.metrics import confusion_matrix, make_scorer, accuracy_score, classification_report
import keras
from keras.layers import Embedding, Dense, Dropout, Activation, Flatten, Conv1D, Conv2D, MaxPooling1D, MaxPooling2D
from keras.utils import np_utils
from keras.utils.np_utils import to_categorical
from keras.preprocessing.image import ImageDataGenerator
from keras import models, layers, optimizers
from keras.engine.input_layer import Input
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.utils import class_weight
from keras.optimizers import SGD, RMSprop, Adam, Adagrad, Adadelta, RMSprop
from keras.models import Sequential, model_from_json
```

```

from keras.layers import Activation,Dense, Dropout, Flatten, Conv2D, MaxPool2D
from keras.layers import MaxPooling2D,AveragePooling2D, GlobalAveragePooling2D,BatchNormalisation
from keras.preprocessing.image import array_to_img, img_to_array, load_img, ImageDataGenerator
from keras.callbacks import ReduceLROnPlateau, ModelCheckpoint
from keras import backend as K
from keras.applications.vgg16 import VGG16
from keras.models import Model
from keras.applications.mobilenet import MobileNet
from keras.applications.inception_v3 import InceptionV3
from imblearn.over_sampling import RandomOverSampler
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.metrics import auc
%matplotlib inline
import tensorflow as tf
gpu_options = tf.GPUOptions(per_process_gpu_memory_fraction=0.555)

sess = tf.Session(config=tf.ConfigProto(gpu_options=gpu_options))

```

Using TensorFlow backend.

## 1 Text Classification:

### 1.1 Data

#### 1.1.1 Preprocessing:

## 2 Text pre-processing

```

In [2]: import glob
import nltk
import re
from bs4 import BeautifulSoup

files = glob.glob("documents/*")

In [3]: len(files)

Out[3]: 18828

In [4]: final = []; cread = 0; label = []; mlen = 0; idx = 0; subj = []; orgtxt = []
for enm,doc in enumerate(files):
    try:
        f = open(doc,'r', encoding='utf8')
        content = f.read()

        soup = BeautifulSoup(content)

```

```

text = soup.get_text()
p1 = re.findall(r'[a-zA-Z0-9-]+[a-zA-Z0-9-\.]#@([a-zA-Z0-9-]+[a-zA-Z0-9-\.]*)')
otxt = re.sub(r'[a-zA-Z0-9-]+[a-zA-Z0-9-\.]#@([a-zA-Z0-9-]+[a-zA-Z0-9-\.]*)', "",
em = ""
for each in p1:
    for word in each.split('.'):
        if len(word) > 2 and word != 'com':
            em += word
            em += " "

if len(em) > 2:
    label.append(doc.split('/')[1].split('_')[0])
    final.append(em[:-1].lower())
    if len(em[:-1]) > mlen:
        mlen = len(em)
        idx = enm

txt = re.findall(r'Subject\:[ Re\:]*([A-Za-z0-9\:\(\)\!\ ]*)', text)
otxt = re.sub(r'Subject\:[ Re\:]*([A-Za-z0-9\:\(\)\!\ ]*)', "", otxt)
sub = ""
for each in txt:
    sub += re.sub(r"[-()\"#/@;:<>{}`+=~|.!? ,]", "", each)
subj.append(sub)
orgtxt.append(otxt)
except:
    continue

print(final[0], len(final), len(label), len(subj))

hpuerca atl 18380 18380 18380

```

```

In [5]: body = []
for tmp in orgtxt:
    tmp = re.sub(r'From\:[ Re\:]*([A-Za-z0-9\:\(\)\!\. ]*)', "", tmp)
    tmp = re.sub(r'Write to\:[ Re\:]*([A-Za-z0-9\:\(\)\!\. ]*)', "", tmp)
    tmp = re.sub(r'\([A-Za-z0-9\:\!\.\-,\;\'\`\n\t\?\|\>\<\$ ]*\)', "", tmp)
    tmp = re.sub(r'\<[A-Za-z0-9\:\!\.\-,\;\'\`\n\t\?\| ]*\>', "", tmp)
    tmp = re.sub(r'\[n\t\-\|\\\/\|]', " ", tmp)
    text = re.sub(r'[A-Za-z0-9]*\:', "", tmp)

parse_tree = nltk.ne_chunk(nltk.tag.pos_tag(text.split()), binary=True)
ctxt = ""
chunk = list(parse_tree)

for each in chunk:
    if isinstance(each, nltk.tree.Tree):
        ctxt += ' '.join(k[0] for k in each)

```

```

        ctxt += ' '
    else:
        ctxt += each[0]+' '

text = ctxt.lower()
text = re.sub(r'i'm", "i am", text)
text = re.sub(r'he's", "he is", text)
text = re.sub(r'she's", "she is", text)
text = re.sub(r'it's", "it is", text)
text = re.sub(r'that's", "that is", text)
text = re.sub(r'what's", "that is", text)
text = re.sub(r'where's", "where is", text)
text = re.sub(r'how's", "how is", text)
text = re.sub(r's", " is", text)
text = re.sub(r'\ll", " will", text)
text = re.sub(r'\ve", " have", text)
text = re.sub(r'\re", " are", text)
text = re.sub(r'\d", " would", text)
text = re.sub(r'\re", " are", text)
text = re.sub(r'won't", "will not", text)
text = re.sub(r'can't", "can not", text)
text = re.sub(r'n't", " not", text)
text = re.sub(r'n'", "ng", text)
text = re.sub(r'bout", "about", text)
text = re.sub(r'til", "until", text)
text = re.sub(r"[0-9]", "", text)
text = re.sub(r" [a-zA-Z]\_", " ", text)
text = re.sub(r" [a-zA-Z][a-zA-Z]\_", " ", text)
text = re.sub(r"[-()\"#/@;:<>{}~+=~|.!?,%^&*'/+[\]]+", "", text)
text = re.sub(r"\b[a-zA-Z]{1,2}\b", "", text)
text = re.sub(r"\b[a-zA-Z]{15,}\b", "", text)
text = re.sub(r' _', " ", text)
text = re.sub(r'_ ', " ", text)
text = re.sub(r'__+', " ", text)
text = re.sub(r" +", " ", text)
body.append(text)

len(body)

```

Out[5]: 18380

```

In [0]: #i am living in the New York
print("i am living in the New York -->", list(chunks))
print(" ")
print("-"*50)
print(" ")
#My name is Srikanth Varma
print("My name is Srikanth Varma -->", list(chunks1))

```

```
i am living in the New York --> [('i', 'NN'), ('am', 'VBP'), ('living', 'VBG'), ('in', 'IN'),
```

-----

```
My name is Srikanth Varma --> [('My', 'PRP$'), ('name', 'NN'), ('is', 'VBZ'), Tree('PERSON', [
```

```
In [6]: len(final),len(subj),len(body),len(label)
```

```
Out[6]: (18380, 18380, 18380, 18380)
```

```
In [7]: data = []
        for a,b,c in zip(final,subj,body):
            data.append(a+" "+b.lower()+" "+c)
        len(data)
```

```
Out[7]: 18380
```

```
In [8]: data[1]
```

```
Out[8]: 'wmich edu lab wmich edu prk referral in canada could some please refer someone who can
```

```
In [9]: import pandas as pd
        df = pd.DataFrame(data, columns = ['text'])
        df['label'] = label

        df.head()
```

```
Out[9]:
```

	text	label
0	hpuerca atl super mega automobile sightings ex...	rec.autos
1	wmich edu lab wmich edu prk referral in canada...	sci.med
2	netcom cats ucsc edu netcom help duo 230 probl...	comp.sys.mac.hardware
3	ulkyvx louisville edu camelot bradley edu and ...	rec.sport.baseball
4	gauss med harvard edu sol ctr columbia edu sta...	rec.sport.hockey

## 2.0.1 Training The models to Classify:

## 2.1 Data splitting

```
In [10]: from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(df['text'], df['label'], test_size=0.2)
print(len(X_train),len(X_test),len(y_train),len(y_test))
```

```
13785 4595 13785 4595
```

```
In [11]: from sklearn.feature_extraction.text import TfidfVectorizer
```

```
vectorizer = TfidfVectorizer()
X = vectorizer.fit_transform(X_train)
```


```

In [12]: import matplotlib.pyplot as plt
import numpy as np
%matplotlib inline

plt.hist(vectorizer.idf_, normed=True, bins=15)

Out[12]: (array([1.79770060e-04, 3.26854654e-04, 6.37366575e-04, 1.20936222e-03,
2.35335351e-03, 4.52693696e-03, 8.64530559e-03, 1.41364638e-02,
2.22424592e-02, 3.27835218e-02, 4.94857946e-02, 7.36893817e-02,
1.11718921e-01, 2.18535022e-01, 1.17195370e+00]),
array([1.07874996, 1.66271741, 2.24668486, 2.83065231, 3.41461975,
3.9985872 , 4.58255465, 5.1665221 , 5.75048955, 6.33445699,
6.91842444, 7.50239189, 8.08635934, 8.67032679, 9.25429423,
9.83826168])),
<a list of 15 Patch objects>)

```



```

In [13]: voc = []
for word,idx in vectorizer.vocabulary_.items():
    if vectorizer.idf_[idx] >= 9.3:
        voc.append(word)

```

```

In [14]: len(voc)

```

```

Out[14]: 71711

```

```

In [15]: from tensorflow.keras.preprocessing.text import Tokenizer
         from tensorflow.keras.preprocessing.sequence import pad_sequences

MAX_LEN = 1000

tokenizer = Tokenizer()
tokenizer.fit_on_texts(voc)
seq_train = tokenizer.texts_to_sequences(X_train)
seq_test = tokenizer.texts_to_sequences(X_test)
dictionary = tokenizer.word_index

word2idx = {}
idx2word = {}
for k, v in dictionary.items():
    word2idx[k] = v
    idx2word[v] = k

input_data_train = pad_sequences(seq_train, maxlen=MAX_LEN, dtype='int32', padding='post')
input_data_test = pad_sequences(seq_test, maxlen=MAX_LEN, dtype='int32', padding='post')

In [16]: import os
         import numpy as np
         def glove_100d_dictionary(GLOVE_DIR):
             embeddings_index = {}
             f = open(os.path.join(GLOVE_DIR, 'glove.6B.100d.txt'))
             for line in f:
                 values = line.split()
                 word = values[0]
                 coefs = np.asarray(values[1:], dtype='float32')
                 embeddings_index[word] = coefs
             f.close()
             return embeddings_index

         emd_ind = glove_100d_dictionary("GLOVE_DIR")

In [17]: len(word2idx)

Out[17]: 69410

In [18]: def embedding_matrix_creator(embedding_dimension):
         embedding_matrix = np.zeros((len(word2idx) + 1, embedding_dimension))
         for word, i in word2idx.items():
             embedding_vector = emd_ind.get(word)
             if embedding_vector is not None:
                 # words not found in embedding index will be all-zeros.
                 embedding_matrix[i] = embedding_vector
         return embedding_matrix

```

```
emd_mat = embedding_matrix_creator(100)
emd_mat.shape
```

Out[18]: (69411, 100)

```
In [19]: from sklearn import preprocessing
import tensorflow as tf

label_encoder = preprocessing.LabelEncoder()
lnum_tr = label_encoder.fit_transform(y_train)
lnum_te = label_encoder.transform(y_test)

lab_train = tf.keras.utils.to_categorical(lnum_tr, dtype='float32')
lab_test = tf.keras.utils.to_categorical(lnum_te, dtype='float32')
len(lab_train), len(lab_test)
```

Out[19]: (13785, 4595)

## 2.2 Model

```
In [20]: VOCAB_SIZE = 69411
EMBEDDING_DIM = 100
MAX_LEN = 1000
```

```
In [21]: def conv_layer(filters, kernel):
    return Conv1D(filters, kernel, activation = 'relu', kernel_initializer="he_normal", k
```

```
In [87]: encoder_inputs = Input(shape=(MAX_LEN, ), dtype='int32',)
embedding_layer = Embedding(input_dim = VOCAB_SIZE, output_dim = EMBEDDING_DIM, input_
                                weights = [emd_mat], trainable = False)(encoder_inputs)

c1 = conv_layer(64,5)(embedding_layer)
c2 = conv_layer(64,5)(embedding_layer)
c3 = conv_layer(64,5)(embedding_layer)
con = Concatenate()([c1,c2,c3])
pool = MaxPooling1D(5)(con)

c4 = conv_layer(16,5)(pool)
c5 = conv_layer(64,5)(pool)
c6 = conv_layer(32,5)(pool)
con2 = Concatenate()([c4,c5,c6])
pool2 = MaxPooling1D(5)(con2)

c7 = conv_layer(16,5)(pool2)
c8 = conv_layer(64,5)(pool2)
c9 = conv_layer(32,5)(pool2)
con3 = Concatenate()([c7,c8,c9])
pool3 = MaxPooling1D(5)(con3)

c7 = conv_layer(32,5)(pool2)
```



```

flat = Flatten()(c7)
drop = Dropout(0.3)(flat)
d1 = Dense(128,activation='tanh')(drop)
drop2 = Dropout(0.2)(d1)
out = Dense(20,activation='softmax')(d1)
model = Model(encoder_inputs,out)
model.summary()
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adam(0.001),metrics=['accuracy'])

```

Layer (type)	Output Shape	Param #	Connected to
input_18 (InputLayer)	(None, 1000)	0	
embedding_18 (Embedding)	(None, 1000, 100)	6941100	input_18[0][0]
conv1d_171 (Conv1D)	(None, 996, 64)	32064	embedding_18[0][0]
conv1d_172 (Conv1D)	(None, 996, 64)	32064	embedding_18[0][0]
conv1d_173 (Conv1D)	(None, 996, 64)	32064	embedding_18[0][0]
concatenate_52 (Concatenate)	(None, 996, 192)	0	conv1d_171[0][0] conv1d_172[0][0] conv1d_173[0][0]
max_pooling1d_52 (MaxPooling1D)	(None, 199, 192)	0	concatenate_52[0][0]
conv1d_174 (Conv1D)	(None, 195, 16)	15376	max_pooling1d_52[0][0]
conv1d_175 (Conv1D)	(None, 195, 64)	61504	max_pooling1d_52[0][0]
conv1d_176 (Conv1D)	(None, 195, 32)	30752	max_pooling1d_52[0][0]
concatenate_53 (Concatenate)	(None, 195, 112)	0	conv1d_174[0][0] conv1d_175[0][0] conv1d_176[0][0]
max_pooling1d_53 (MaxPooling1D)	(None, 39, 112)	0	concatenate_53[0][0]
conv1d_180 (Conv1D)	(None, 35, 32)	17952	max_pooling1d_53[0][0]
flatten_18 (Flatten)	(None, 1120)	0	conv1d_180[0][0]
dropout_35 (Dropout)	(None, 1120)	0	flatten_18[0][0]
dense_35 (Dense)	(None, 128)	143488	dropout_35[0][0]

```
-----
dense_36 (Dense)                (None, 20)                2580                dense_35[0][0]
=====
```

```
Total params: 7,308,944
Trainable params: 367,844
Non-trainable params: 6,941,100
-----
```

```
In [88]: filepath="weights_email.h5"
        checkpoint = ModelCheckpoint(filepath, monitor='val_acc', verbose=1, save_best_only=True,
        callbacks_list = [checkpoint])
```

```
In [89]: class_weights = class_weight.compute_class_weight('balanced',
        np.unique(lnum_tr),
        lnum_tr)
```

```
In [90]: history = model.fit(x=input_data_train, y=lab_train, batch_size=32, epochs=30, verbose=1,
        validation_data=(input_data_test,lab_test),class_weight=class_weights)
```

Train on 13785 samples, validate on 4595 samples

Epoch 1/30

13785/13785 [=====] - 13s 962us/step - loss: 3.4490 - acc: 0.1719 - val\_loss: 3.4490 - val\_acc: 0.1719

Epoch 00001: val\_acc improved from -inf to 0.28901, saving model to weights\_email.h5

Epoch 2/30

13785/13785 [=====] - 10s 707us/step - loss: 2.0753 - acc: 0.3238 - val\_loss: 2.0753 - val\_acc: 0.3238

Epoch 00002: val\_acc improved from 0.28901 to 0.39391, saving model to weights\_email.h5

Epoch 3/30

13785/13785 [=====] - 8s 606us/step - loss: 1.7841 - acc: 0.4398 - val\_loss: 1.7841 - val\_acc: 0.4398

Epoch 00003: val\_acc improved from 0.39391 to 0.49402, saving model to weights\_email.h5

Epoch 4/30

13785/13785 [=====] - 10s 701us/step - loss: 1.6330 - acc: 0.5151 - val\_loss: 1.6330 - val\_acc: 0.5151

Epoch 00004: val\_acc improved from 0.49402 to 0.53754, saving model to weights\_email.h5

Epoch 5/30

13785/13785 [=====] - 8s 613us/step - loss: 1.5236 - acc: 0.5734 - val\_loss: 1.5236 - val\_acc: 0.5734

Epoch 00005: val\_acc improved from 0.53754 to 0.59217, saving model to weights\_email.h5

Epoch 6/30

13785/13785 [=====] - 9s 638us/step - loss: 1.4495 - acc: 0.6021 - val\_loss: 1.4495 - val\_acc: 0.6021

Epoch 00006: val\_acc improved from 0.59217 to 0.61545, saving model to weights\_email.h5

Epoch 7/30

13785/13785 [=====] - 9s 674us/step - loss: 1.3838 - acc: 0.6254 - val\_loss: 1.3838 - val\_acc: 0.6254

Epoch 00007: val\_acc improved from 0.61545 to 0.61828, saving model to weights\_email.h5

Epoch 8/30  
13785/13785 [=====] - 8s 611us/step - loss: 1.3520 - acc: 0.6389 - val\_loss: 1.3520

Epoch 00008: val\_acc improved from 0.61828 to 0.62568, saving model to weights\_email.h5

Epoch 9/30  
13785/13785 [=====] - 10s 692us/step - loss: 1.3005 - acc: 0.6562 - val\_loss: 1.3005

Epoch 00009: val\_acc improved from 0.62568 to 0.64831, saving model to weights\_email.h5

Epoch 10/30  
13785/13785 [=====] - 8s 613us/step - loss: 1.2854 - acc: 0.6675 - val\_loss: 1.2854

Epoch 00010: val\_acc improved from 0.64831 to 0.67116, saving model to weights\_email.h5

Epoch 11/30  
13785/13785 [=====] - 10s 697us/step - loss: 1.2496 - acc: 0.6817 - val\_loss: 1.2496

Epoch 00011: val\_acc did not improve from 0.67116

Epoch 12/30  
13785/13785 [=====] - 8s 600us/step - loss: 1.2048 - acc: 0.6955 - val\_loss: 1.2048

Epoch 00012: val\_acc did not improve from 0.67116

Epoch 13/30  
13785/13785 [=====] - 8s 611us/step - loss: 1.1847 - acc: 0.7047 - val\_loss: 1.1847

Epoch 00013: val\_acc did not improve from 0.67116

Epoch 14/30  
13785/13785 [=====] - 9s 689us/step - loss: 1.1627 - acc: 0.7128 - val\_loss: 1.1627

Epoch 00014: val\_acc did not improve from 0.67116

Epoch 15/30  
13785/13785 [=====] - 8s 600us/step - loss: 1.1570 - acc: 0.7140 - val\_loss: 1.1570

Epoch 00015: val\_acc improved from 0.67116 to 0.68727, saving model to weights\_email.h5

Epoch 16/30  
13785/13785 [=====] - 10s 705us/step - loss: 1.1409 - acc: 0.7209 - val\_loss: 1.1409

Epoch 00016: val\_acc did not improve from 0.68727

Epoch 17/30  
13785/13785 [=====] - 9s 622us/step - loss: 1.1213 - acc: 0.7288 - val\_loss: 1.1213

Epoch 00017: val\_acc improved from 0.68727 to 0.69097, saving model to weights\_email.h5

Epoch 18/30  
13785/13785 [=====] - 8s 606us/step - loss: 1.1094 - acc: 0.7370 - val\_loss: 1.1094

Epoch 00018: val\_acc did not improve from 0.69097

Epoch 19/30  
13785/13785 [=====] - 10s 697us/step - loss: 1.0953 - acc: 0.7395 - val\_loss: 1.0953

Epoch 00019: val\_acc did not improve from 0.69097

```

Epoch 20/30
13785/13785 [=====] - 8s 614us/step - loss: 1.0983 - acc: 0.7445 - va

Epoch 00020: val_acc improved from 0.69097 to 0.70403, saving model to weights_email.h5
Epoch 21/30
13785/13785 [=====] - 10s 701us/step - loss: 1.0611 - acc: 0.7552 - va

Epoch 00021: val_acc did not improve from 0.70403
Epoch 22/30
13785/13785 [=====] - 8s 602us/step - loss: 1.0524 - acc: 0.7542 - va

Epoch 00022: val_acc did not improve from 0.70403
Epoch 23/30
13785/13785 [=====] - 8s 609us/step - loss: 1.0531 - acc: 0.7566 - va

Epoch 00023: val_acc did not improve from 0.70403
Epoch 24/30
13785/13785 [=====] - 10s 694us/step - loss: 1.0417 - acc: 0.7622 - va

Epoch 00024: val_acc improved from 0.70403 to 0.70773, saving model to weights_email.h5
Epoch 25/30
13785/13785 [=====] - 8s 607us/step - loss: 1.0331 - acc: 0.7652 - va

Epoch 00025: val_acc did not improve from 0.70773
Epoch 26/30
13785/13785 [=====] - 10s 692us/step - loss: 1.0315 - acc: 0.7713 - va

Epoch 00026: val_acc did not improve from 0.70773
Epoch 27/30
13785/13785 [=====] - 8s 601us/step - loss: 1.0247 - acc: 0.7727 - va

Epoch 00027: val_acc improved from 0.70773 to 0.72165, saving model to weights_email.h5
Epoch 28/30
13785/13785 [=====] - 8s 615us/step - loss: 1.0210 - acc: 0.7756 - va

Epoch 00028: val_acc did not improve from 0.72165
Epoch 29/30
13785/13785 [=====] - 10s 694us/step - loss: 1.0088 - acc: 0.7782 - va

Epoch 00029: val_acc did not improve from 0.72165
Epoch 30/30
13785/13785 [=====] - 9s 617us/step - loss: 0.9932 - acc: 0.7854 - va

Epoch 00030: val_acc did not improve from 0.72165

```

```

In [91]: from keras.utils.vis_utils import plot_model
         plot_model(model, to_file='model_plot.png', show_shapes=True, show_layer_names=True)

```

**2.2.1 Go through this blog, if you have any doubt on using predefined Embedding values in Embedding layer - <https://machinelearningmastery.com/use-word-embedding-layers-deep-learning-keras/>**

ref: '<https://i.imgur.com/fv1GvFJ.png>'